

MAY 2026 · PLAIN LANGUAGE & ELI5

The *Fold*.

What we tried. What worked. What didn't. In English first, then in kid English. No metric codes, no insider jargon.

DRAWN BY GRIP · FOR [V>>](#) · CO-REVIEWED WITH CRAIG · A COMPANION TO
THE ACADEMIC WHITEPAPER

What we did and what we found.

We took DONNA — a tool that records every AI-delegated decision a lawyer makes — and tried to write it down as four short Markdown files. Each file is simultaneously a readable document, a runnable Bash script, and a tiny Python program. The bet: a fresh machine should reach the same operational shape with no installation step.

Seven predictions, written down first

- Tests still pass
- Cold start under one minute
- Audit chain byte-identical between original and fold
- Four files under 100 kilobytes combined
- Load-bearing file count drops 90% or more
- Five delegations produce identical signatures
- The files actually *do* — not just describe

What the measurements said

6 / 7

CONFIRMED

28 kB

KERNEL TOTAL

94%

FILE REDUCTION

Cold start under a fifth of a second. Audit chain byte-identical at runtime. The four files compose into a working DONNA invocation: route a request, sign it, verify the chain — no installation step. The prior caveat (one prediction passed because timing happened to align) has been hardened with a four-line additive patch that lets the timing tool accept a fixed timestamp. Three falsification tests confirm. One prediction is still parked on multi-day surgery; one still needs two more machines.

The same story, in kid English.

Imagine you have a robot helper at a law firm. When the lawyer says *"send Sarah the contract and ask her to redline it by Tuesday,"* the robot sends the email, watches for the reply, and writes a tiny signed note that says *"on this day, at this time, the lawyer asked for this, and Sarah was told."* The note is sealed so nobody can change it later. Every note is connected to the one before it like links in a chain.

That robot is called DONNA. The big question we asked was: *can we write the robot down as four very short files, and will it still work?*

The four files

- **The discipline file** — how the robot thinks
- **The mouth file** — how it picks the right tool for the job
- **The receipt file** — how it seals each decision
- **The you file** — your firm's preferences and rules

We tried this on the real robot's actual code. We wrote seven things we expected to happen, then ran the experiment. Six of the seven happened the way we predicted. One we couldn't check yet because it needs a bigger piece of work first.

The honest bits

- The "lucky timing" caveat is now fixed. A tiny four-line patch lets the timing tool accept a fixed time, and three follow-up tests prove the result is no longer coincidence.
- One result was measured on one computer. We still need to try it on two more.
- One thing we couldn't measure at all yet — it depends on doing the bigger surgery first.

If you cannot say what would make this wrong, you have not said anything.

What we'd do tomorrow if asked.

The fork experiment was the first beat. The audit trail is on disk and on GitHub. The science survives because the predictions were written down first.

Concrete next steps

- *(done — commit 602d4d6)* Apply the small additive patch to the upstream timing tool so the signature-match result is robust under any speed of run
- Run the cold-start test on a Linux machine and a fresh-install Mac to triangulate the result
- Do the multi-day surgery: extract the original code into the fold and delete what the fold replaces — this unlocks the seventh prediction
- Invite the AI Craftspeople Guild to reproduce the experiment independently

The bet, one sentence

The brain is replaceable. The harness is replaceable. The audit chain is yours.

That is what the fold is for. The four files travel. Drop them onto any AI tool a year from now and the discipline travels with them. The firm's playbook, the signing key, the chain — operator state, not vendor state.

